



Name : P. Manikandan

Reg no: 814421243014

Dept : Artificial intelligence and data science

Title : Sentimental Analysis for Marketing




Table of contents:

- Executive Summary
- Introduction
- Problem Statement Revisited
- Data Collection and Preprocessing
- Sentiment Analysis Techniques
- Feature Extraction
- Sentiment Analysis Results
- Advanced Technique
- Business Insights
- Conclusion
- Acknowledgments

1. Executive Summary:

The Executive Summary provides a concise yet comprehensive overview of the entire project. It should include:

- This phase involves summarizing the project's journey, methodologies, and results, ensuring that no important detail is missed.
- It represents the closure of the analysis process and the beginning of the implementation of valuable insights in real-world business decisions.
- The project documentation serves as a lasting record, offering reference material for future work and presentations to stakeholders.

2. Introduction:

- The project focuses on sentiment analysis and its significance in understanding customer perceptions of competitor products.
- Understanding customer sentiments is vital for improving products and enhancing marketing strategies.
- This project utilizes various NLP methods to extract insights from customer feedback.
- It encompasses problem definition, data collection, preprocessing, advanced sentiment analysis, and visualization.
- The project aims to enhance sentiment analysis accuracy by exploring advanced techniques.
- The ultimate goal is to generate actionable business insights.

3. Problem Statement Revisited:

- The problem is to perform sentiment analysis on customer feedback to gain insights into competitor products."

4. Data Collection and Preprocessing:

- This section should provide detailed information on data collection and preprocessing, including:

Data Collection:

- The project begins with the essential step of data collection. A dataset containing customer reviews and sentiments about competitor products is identified for analysis.
- The provided dataset comes from a reputable source, such as Kaggle, and contains relevant data for sentiment analysis.

Data Preprocessing:

- Data preprocessing is a critical phase in the project, involving cleaning and transforming the textual data for analysis.
- Textual data is carefully cleaned to remove noise and irrelevant information, ensuring that the dataset is suitable for analysis.

5. Sentiment Analysis Techniques:

- The project utilizes various Natural Language Processing (NLP) methods, including Bag of Words (BoW), Word Embeddings, and Transformer models like BERT and RoBERTa.

- These techniques enable the analysis of customer sentiments and opinions expressed in the textual data.
- Specific NLP libraries and tools, such as NLTK, spaCy, and Hugging Face Transformers, are employed to implement these techniques effectively.

6. Feature Extraction:

- Feature extraction is a critical step in your project, where textual data is transformed into numerical representations suitable for sentiment analysis.
- This phase is responsible for converting raw text into structured features that can be analyzed effectively, facilitating the understanding of customer sentiments.
- The chosen techniques and algorithms for feature extraction play a pivotal role in revealing the underlying patterns and nuances within the data.

7. Sentiment Analysis Results:

- Sentiment analysis results form a core part of the project, presenting insights into customer feedback and opinions.
- These results showcase the distribution of sentiments, trends in customer responses, and key findings that shed light on the strengths and weaknesses of competitor products.
- The use of visualizations, such as charts and graphs, aids in presenting the sentiment analysis findings in a clear and accessible manner.
- The analysis results help in drawing actionable business insights, guiding marketing strategies, product improvements, and competition analysis.

8. Advanced Techniques:

- Advanced techniques represent a pivotal phase in your project, where fine-tuning pre-trained sentiment analysis models, such as BERT and RoBERTa, are investigated for improving prediction accuracy.
- These state-of-the-art models are leveraged to capture nuanced sentiments and contextual understanding within customer feedback.
- The application of advanced techniques signifies the project's commitment to enhancing the quality of sentiment predictions.

9. Business Insights:

Summarize the valuable insights generated from the sentiment analysis results:

- Business insights are the heart of the project, offering actionable guidance to shape marketing strategies and product development.
- These insights are distilled from the sentiment analysis results and highlight areas of improvement, customer preferences, and competitive strengths and weaknesses.
- They are a valuable asset for making data-driven decisions in the ever-evolving business landscape.

10. Conclusion:

- The conclusion marks the culmination of the project, summarizing the key findings and their implications for marketing and product development.
- It highlights the significance of sentiment analysis in the modern business landscape, emphasizing the value of understanding customer feedback.

- The insights obtained from the sentiment analysis results contribute to data-driven decision-making, providing a roadmap for improvement and competition analysis.
- The project underlines the power of NLP techniques in transforming raw textual data into actionable insights.
- The journey from problem definition to insight generation is a testament to the project's ability to enhance marketing strategies and guide businesses toward success in a competitive market.

11. Acknowledgments:

- Acknowledge individuals or organizations that provided support, data, or resources during the project:
- Express gratitude for their contributions or assistance.

Program:

```
[1]: import pandas as pd
import numpy as np
# %load_ext nb_black

# library to suppress warnings or deprecation notes
import warnings

warnings.filterwarnings("ignore")

# import Regex, string and unicodedata.
import re, string, unicodedata

import contractions

# import BeautifulSoup.
from bs4 import BeautifulSoup

# import Natural Language Tool-Kit.
import nltk

# download Stopwords.
nltk.download("stopwords")
nltk.download("punkt")
nltk.download("wordnet")

# import stopwords.
from nltk.corpus import stopwords

# import Tokenizer.
from nltk.tokenize import word_tokenize, sent_tokenize

# library to split data
from sklearn.model_selection import train_test_split, StratifiedKFold

# libraries to help with data visualization
import matplotlib.pyplot as plt
```



```

import seaborn as sns
import missingno as msno

# import wordcloud
import wordcloud
from wordcloud import STOPWORDS
from wordcloud import WordCloud

# remove the limit for the number of displayed columns
pd.set_option("display.max_columns", None)

# set the limit for the number of displayed rows
pd.set_option("display.max_rows", 200)

# to get different metric scores
from sklearn.metrics import (
    recall_score,
    accuracy_score,
    confusion_matrix, classification_report,
    f1_score,
    precision_score,
    precision_recall_fscore_support
)

# import vectorizers
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# import rfc and cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score

# import word preprocessors
from nltk.tokenize import word_tokenize
from nltk.stem import LancasterStemmer, WordNetLemmatizer

```

```

[nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\Administrator\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\Administrator\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data]   C:\Users\Administrator\AppData\Roaming\nltk_data...
[nltk_data]   Package wordnet is already up-to-date!

```

```

[2]: df = pd.read_csv("Tweets.csv")
      df.head()

```

```
[2]:      tweet_id  airline_sentiment  airline_sentiment_confidence \
0  570306133677760513          neutral          1.0000
1  570301130888122368          positive          0.3486
2  570301083672813571          neutral          0.6837
3  570301031407624196          negative          1.0000
4  570300817074462722          negative          1.0000

      negativereason  negativereason_confidence      airline \
0              NaN              NaN  Virgin America
1              NaN              0.0000  Virgin America
2              NaN              NaN  Virgin America
3      Bad Flight              0.7033  Virgin America
4      Can't Tell              1.0000  Virgin America

      airline_sentiment_gold      name  negativereason_gold  retweet_count \
0              NaN      cairdin              NaN              0
1              NaN      jnardino              NaN              0
2              NaN  yvonnalynn              NaN              0
3              NaN      jnardino              NaN              0
4              NaN      jnardino              NaN              0

      text  tweet_coord \
0      @VirginAmerica What @dhepburn said.              NaN
1      @VirginAmerica plus you've added commercials t...              NaN
2      @VirginAmerica I didn't today... Must mean I n...              NaN
3      @VirginAmerica it's really aggressive to blast...              NaN
4      @VirginAmerica and it's a really big bad thing...              NaN

      tweet_created  tweet_location      user_timezone
0  2015-02-24 11:35:52 -0800              NaN  Eastern Time (US & Canada)
1  2015-02-24 11:15:59 -0800              NaN  Pacific Time (US & Canada)
2  2015-02-24 11:15:48 -0800      Lets Play  Central Time (US & Canada)
3  2015-02-24 11:15:36 -0800              NaN  Pacific Time (US & Canada)
4  2015-02-24 11:14:45 -0800              NaN  Pacific Time (US & Canada)
```

```
[3]: texts = [[word.lower() for word in text.split()] for text in df]
df.head()
```

```
[3]:      tweet_id  airline_sentiment  airline_sentiment_confidence \
0  570306133677760513          neutral          1.0000
1  570301130888122368          positive          0.3486
2  570301083672813571          neutral          0.6837
3  570301031407624196          negative          1.0000
4  570300817074462722          negative          1.0000

      negativereason  negativereason_confidence      airline \
0              NaN              NaN  Virgin America
```

1	NaN	0.0000	Virgin America
2	NaN	NaN	Virgin America
3	Bad Flight	0.7033	Virgin America
4	Can't Tell	1.0000	Virgin America

	airline_sentiment_gold	name	negativereason_gold	retweet_count	\
0	NaN	cairdin	NaN	0	
1	NaN	jnardino	NaN	0	
2	NaN	yvonnalynn	NaN	0	
3	NaN	jnardino	NaN	0	
4	NaN	jnardino	NaN	0	

	text	tweet_coord	\
0	@VirginAmerica What @dhepburn said.	NaN	
1	@VirginAmerica plus you've added commercials t...	NaN	
2	@VirginAmerica I didn't today... Must mean I n...	NaN	
3	@VirginAmerica it's really aggressive to blast...	NaN	
4	@VirginAmerica and it's a really big bad thing...	NaN	

	tweet_created	tweet_location	user_timezone
0	2015-02-24 11:35:52 -0800	NaN	Eastern Time (US & Canada)
1	2015-02-24 11:15:59 -0800	NaN	Pacific Time (US & Canada)
2	2015-02-24 11:15:48 -0800	Lets Play	Central Time (US & Canada)
3	2015-02-24 11:15:36 -0800	NaN	Pacific Time (US & Canada)
4	2015-02-24 11:14:45 -0800	NaN	Pacific Time (US & Canada)

[4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 14640 entries, 0 to 14639
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	tweet_id	14640 non-null	int64
1	airline_sentiment	14640 non-null	object
2	airline_sentiment_confidence	14640 non-null	float64
3	negativereason	9178 non-null	object
4	negativereason_confidence	10522 non-null	float64
5	airline	14640 non-null	object
6	airline_sentiment_gold	40 non-null	object
7	name	14640 non-null	object
8	negativereason_gold	32 non-null	object
9	retweet_count	14640 non-null	int64
10	text	14640 non-null	object
11	tweet_coord	1019 non-null	object
12	tweet_created	14640 non-null	object
13	tweet_location	9907 non-null	object

```
14 user_timezone          9820 non-null object
dtypes: float64(2), int64(2), object(11)
memory usage: 1.7+ MB
```

```
[5]: df.isnull().sum()
```


```
[5]: tweet_id          0
     airline_sentiment  0
     airline_sentiment_confidence  0
     negativereason     5462
     negativereason_confidence  4118
     airline            0
     airline_sentiment_gold  14600
     name              0
     negativereason_gold  14608
     retweet_count      0
     text              0
     tweet_coord        13621
     tweet_created       0
     tweet_location     4733
     user_timezone      4820
     dtype: int64
```

```
[6]: df.isnull().sum() / len(df) * 100
```

```
[6]: tweet_id          0.000000
     airline_sentiment  0.000000
     airline_sentiment_confidence  0.000000
     negativereason     37.308743
     negativereason_confidence  28.128415
     airline            0.000000
     airline_sentiment_gold  99.726776
     name              0.000000
     negativereason_gold  99.781421
     retweet_count      0.000000
     text              0.000000
     tweet_coord        93.039617
     tweet_created       0.000000
     tweet_location     32.329235
     user_timezone      32.923497
     dtype: float64
```


```
[7]: msno.matrix(df)
```

```
[7]: <AxesSubplot:>
```



```
[8]: plt.figure(figsize=(12,7))
sns.heatmap(df.isnull(), cmap = "Blues")
      ↳ of missing value using heatmap
plt.title("Missing values?", fontsize = 15)
plt.show()
```

#Visualization_



```
[9]: print("Percentage null or na values in df")
      ((df.isnull() | df.isna()).sum() * 100 / df.index.size).round(2)
```

Percentage null or na values in df

```
[9]: tweet_id          0.00
      airline_sentiment  0.00
      airline_sentiment_confidence  0.00
      negativereason    37.31
      negativereason_confidence  28.13
      airline           0.00
      airline_sentiment_gold  99.73
      name              0.00
      negativereason_gold  99.78
      retweet_count     0.00
      text              0.00
      tweet_coord       93.04
```

```

tweet_created          0.00
tweet_location         32.33
user_timezone          32.92
dtype: float64

```

```

[10]: del df["tweet_coord"]
      del df["airline_sentiment_gold"]
      del df["negativereason_gold"]

```

```

[11]: df.head()

```

```

[11]:      tweet_id  airline_sentiment  airline_sentiment_confidence \
0  570306133677760513          neutral                1.0000
1  570301130888122368         positive                0.3486
2  570301083672813571          neutral                0.6837
3  570301031407624196         negative                1.0000
4  570300817074462722         negative                1.0000

      negativereason  negativereason_confidence  airline  name \
0              NaN                NaN  Virgin America  cairdin
1              NaN                0.0000  Virgin America  jnardino
2              NaN                NaN  Virgin America  yvonnalynn
3      Bad Flight                0.7033  Virgin America  jnardino
4      Can't Tell                1.0000  Virgin America  jnardino

      retweet_count  text \
0              0  @VirginAmerica What @dhepburn said.
1              0  @VirginAmerica plus you've added commercials t...
2              0  @VirginAmerica I didn't today... Must mean I n...
3              0  @VirginAmerica it's really aggressive to blast...
4              0  @VirginAmerica and it's a really big bad thing...

      tweet_created  tweet_location  user_timezone
0  2015-02-24 11:35:52 -0800      NaN  Eastern Time (US & Canada)
1  2015-02-24 11:15:59 -0800      NaN  Pacific Time (US & Canada)
2  2015-02-24 11:15:48 -0800  Lets Play  Central Time (US & Canada)
3  2015-02-24 11:15:36 -0800      NaN  Pacific Time (US & Canada)
4  2015-02-24 11:14:45 -0800      NaN  Pacific Time (US & Canada)

```

```

[12]: freq = df.groupby("negativereason").size()

```

```

[13]: # Checking duplicates
      df.duplicated().sum()

```

```

[13]: 39

```

```
[14]: df.drop_duplicates(inplace = True)
df.duplicated().sum()
```

```
[14]: 0
```

```
[15]: df.sample(n = 10)
```

```
[15]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	\
10589	569156425626329089	neutral	1.0000	
6182	568149878095753216	neutral	0.6545	
11336	568196165780578304	negative	1.0000	
623	570245555064074240	negative	1.0000	
1186	569902065247322112	negative	1.0000	
2425	569213883371683840	positive	0.6679	
13299	569893723091238912	negative	1.0000	
7693	569343003476819969	neutral	0.6641	
5148	569308552671707136	negative	1.0000	
11135	568486436355346432	negative	1.0000	

	negativereason	negativereason_confidence	airline	\
10589	NaN	NaN	US Airways	
6182	NaN	0.0000	Southwest	
11336	Can't Tell	0.3579	US Airways	
623	Flight Booking Problems	0.6740	United	
1186	Late Flight	1.0000	United	
2425	NaN	NaN	United	
13299	longlines	0.3512	American	
7693	NaN	0.0000	Delta	
5148	Lost Luggage	1.0000	Southwest	
11135	Bad Flight	1.0000	US Airways	

	name	retweet_count	\
10589	observepeople	0	
6182	Brian_Fox	0	
11336	thefisch26	0	
623	fatwmnonthemtn	0	
1186	LukeXuanLiu	1	
2425	PierreSchmit	0	
13299	elisakathleen	0	
7693	dgruber1700	0	
5148	scoobydoo9749	0	
11135	kristenlc	0	

	text	\
10589	@usairways Does anyone know the hold times for...	
6182	@SouthwestAir I would but you need to follow m...	
11336	@USAirways Secondary screenings, a piece of th...	


```

623    @united What's going on with your website? I'm...
1186   @united and most frustratingly, all this delay...
2425   @united gave me a smile today, with a Zero Awa...
13299  @AmericanAir the most stressful morning and st...
7693   @JetBlue flite454
5148   @SouthwestAir 9 hrs in Baltimore, still not go...
11135  @USAirways we bought our tickets months ago. H...

```

	tweet_created	tweet_location \
10589	2015-02-21 07:27:20 -0800	NaN
6182	2015-02-18 12:47:41 -0800	NH, United States
11336	2015-02-18 15:51:37 -0800	Washington, DC
623	2015-02-24 07:35:09 -0800	Summit, NJ
1186	2015-02-23 08:50:15 -0800	NaN
2425	2015-02-21 11:15:39 -0800	Rixensart, Belgium
13299	2015-02-23 08:17:06 -0800	Boston, MA
7693	2015-02-21 19:48:44 -0800	NaN
5148	2015-02-21 17:31:50 -0800	Tallahassee, FL
11135	2015-02-19 11:05:03 -0800	NaN

	user_timezone
10589	Eastern Time (US & Canada)
6182	Eastern Time (US & Canada)
11336	Central Time (US & Canada)
623	Central Time (US & Canada)
1186	Atlantic Time (Canada)
2425	Brussels
13299	NaN
7693	NaN
5148	America/Chicago
11135	Eastern Time (US & Canada)

```
[16]: df.describe().T
```

```

[16]:
count      mean      std \
tweet_id    14601.0  5.692156e+17  7.782706e+14
airline_sentiment_confidence  14601.0  8.999022e-01  1.629654e-01
negativereason_confidence    10501.0  6.375749e-01  3.303735e-01
retweet_count    14601.0  8.280255e-02  7.467231e-01

min      25%      50% \
tweet_id  5.675883e+17  5.685581e+17  5.694720e+17
airline_sentiment_confidence  3.350000e-01  6.923000e-01  1.000000e+00
negativereason_confidence    0.000000e+00  3.605000e-01  6.705000e-01
retweet_count    0.000000e+00  0.000000e+00  0.000000e+00

75%      max


```

tweet_id	5.698884e+17	5.703106e+17
airline_sentiment_confidence	1.000000e+00	1.000000e+00
negativereason_confidence	1.000000e+00	1.000000e+00
retweet_count	0.000000e+00	4.400000e+01


```
[17]: df.nunique()
```

```
[17]: tweet_id          14485
      airline_sentiment      3
      airline_sentiment_confidence  1023
      negativereason         10
      negativereason_confidence  1410
      airline                6
      name                   7701
      retweet_count          18
      text                   14427
      tweet_created          14247
      tweet_location         3081
      user_timezone          85
      dtype: int64
```

```
[18]: ax = sns.countplot(x = "negativereason_confidence", data = df)
```



```
[19]: plt.figure(figsize = (10, 10))
ax = sns.countplot(x = "airline", data = df)
```



```
[20]: import plotly.graph_objects as go
crosstab_sentiments=pd.crosstab(df.airline, df.negativereason)
companies=list(crosstab_sentiments.index)

fig = go.Figure(data=[
    go.Bar(name=col_name, x=companies, y=list(crosstab_sentiments[col_name]))
for col_name in list(crosstab_sentiments.columns)])#
Change the bar mode
fig.update_layout(barmode='stack',
                    title='Sentiment distribution per company',
                    yaxis=dict(title='Sentiment distribution'));
```

```

xaxis=dict(title="Companies"))
fig.show()

```

```

[21]: crosstab_neg_reasons = pd.crosstab(df["airline"], df["negativereason"])
      companies = list(crosstab_neg_reasons.index)

      fig = go.Figure(data = [
          go.Bar(name = col_name, x = companies, y =
              list(crosstab_neg_reasons[col_name]))
          for col_name in list(crosstab_neg_reasons.columns)])

      fig.update_layout(barmode = "stack",
          title = "Negative Reasons Distribution per Company",
          yaxis = dict(title = "Negative reasons Distribution"),
          xaxis = dict(title = "Companies"))

      fig.show()

```

```

[22]: labels = list(crosstab_neg_reasons.columns)
      values = [crosstab_neg_reasons[col_name].sum() for col_name in labels]

      # Use `hole` to create a donut-like pie chart
      fig = go.Figure(data=[go.Pie(labels=labels, values=values, hole=.3)])
      fig.update_layout(title="Overall distribution for negative reasons")
      fig.show()

```

```

[23]: df.drop(df.loc[df["airline_sentiment"] == "neutral"].index, inplace = True)

```

```

[24]: data = df[["airline_sentiment", "text"]]
      data.head()

```

```

[24]: airline_sentiment      text
1      positive  @VirginAmerica plus you've added commercials t...
3      negative  @VirginAmerica it's really aggressive to blast...
4      negative  @VirginAmerica and it's a really big bad thing...
5      negative  @VirginAmerica seriously would pay $30 a fligh...
6      positive  @VirginAmerica yes, nearly every time I fly VX...

```

```

[25]: X = df["text"]
      y = df["airline_sentiment"]
      X

```

```

[25]: 1      @VirginAmerica plus you've added commercials t...
      3      @VirginAmerica it's really aggressive to blast...
      4      @VirginAmerica and it's a really big bad thing...
      5      @VirginAmerica seriously would pay $30 a fligh...
      6      @VirginAmerica yes, nearly every time I fly VX...

```

```

14633 @AmericanAir my flight was Cancelled Flightled...
14634 @AmericanAir right on cue with the delays
14635 @AmericanAir thank you we got on a different f...
14636 @AmericanAir leaving over 20 minutes Late Flig...
14638 @AmericanAir you have my money, you change my ...
Name: text, Length: 11510, dtype: object

```

```
[26]: y
```

```

[26]: 1      positive
      3      negative
      4      negative
      5      negative
      6      positive
      ...
14633 negative
14634 negative
14635 positive
14636 negative
14638 negative
Name: airline_sentiment, Length: 11510, dtype: object

```

```

[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
random_state = 42)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

```

```
(9208,) (2302,) (9208,) (2302,)
```

```
[28]: tfidf = TfidfVectorizer(stop_words="english")
```

```
[29]: tfidf.fit(y_train)
```

```
[29]: TfidfVectorizer(stop_words='english')
```

```
[30]: print(tfidf.get_feature_names_out())
```

```
['negative' 'positive']
```

```
[31]: print(tfidf.vocabulary_)
```

```
{'negative': 0, 'positive': 1}
```

```
[32]: print(df)
```

	tweet_id	airline_sentiment	airline_sentiment_confidence \
1	570301130888122368	positive	0.3486
3	570301031407624196	negative	1.0000
4	570300817074462722	negative	1.0000

5	570300767074181121	negative	1.0000
6	570300616901320704	positive	0.6745
...
14633	569587705937600512	negative	1.0000
14634	569587691626622976	negative	0.6684
14635	569587686496825344	positive	0.3487
14636	569587371693355008	negative	1.0000
14638	569587188687634433	negative	1.0000

	negativereason	negativereason_confidence	airline \
1	NaN	0.0000	Virgin America
3	Bad Flight	0.7033	Virgin America
4	Can't Tell	1.0000	Virgin America
5	Can't Tell	0.6842	Virgin America
6	NaN	0.0000	Virgin America
...
14633	Cancelled Flight	1.0000	American
14634	Late Flight	0.6684	American
14635	NaN	0.0000	American
14636	Customer Service Issue	1.0000	American
14638	Customer Service Issue	0.6659	American

	name	retweet_count \
1	jnardino	0
3	jnardino	0
4	jnardino	0
5	jnardino	0
6	cjmcginnis	0
...
14633	RussellsWriting	0
14634	GolfWithWoody	0
14635	KristenReenders	0
14636	itsropes	0
14638	SraJackson	0

	text \
1	@VirginAmerica plus you've added commercials t...
3	@VirginAmerica it's really aggressive to blast...
4	@VirginAmerica and it's a really big bad thing...
5	@VirginAmerica seriously would pay \$30 a fligh...
6	@VirginAmerica yes, nearly every time I fly VX...
...	...
14633	@AmericanAir my flight was Cancelled Flightled...
14634	@AmericanAir right on cue with the delays
14635	@AmericanAir thank you we got on a different f...
14636	@AmericanAir leaving over 20 minutes Late Flig...
14638	@AmericanAir you have my money, you change my ...

		tweet_created	tweet_location	user_timezone
1	2015-02-24 11:15:59 -0800	NaN	Pacific Time (US & Canada)	
3	2015-02-24 11:15:36 -0800	NaN	Pacific Time (US & Canada)	
4	2015-02-24 11:14:45 -0800	NaN	Pacific Time (US & Canada)	
5	2015-02-24 11:14:33 -0800	NaN	Pacific Time (US & Canada)	
6	2015-02-24 11:13:57 -0800	San Francisco CA	Pacific Time (US & Canada)	
...	
14633	2015-02-22 12:01:06 -0800	Los Angeles	Arizona	
14634	2015-02-22 12:01:02 -0800	NaN	Quito	
14635	2015-02-22 12:01:01 -0800	NaN	NaN	
14636	2015-02-22 11:59:46 -0800	Texas	NaN	
14638	2015-02-22 11:59:02 -0800	New Jersey	Eastern Time (US & Canada)	

[11510 rows x 12 columns]

```
[33]: data[data["airline_sentiment"] == "negative"]["text"]
```

```
[33]: 3      @VirginAmerica it's really aggressive to blast...
      4      @VirginAmerica and it's a really big bad thing...
      5      @VirginAmerica seriously would pay $30 a fligh...
      15      @VirginAmerica SFO-PDX schedule is still MIA.
      17      @VirginAmerica I flew from NYC to SFO last we...
```

```

      ...
14631  @AmericanAir thx for nothing on getting us out...
14633  @AmericanAir my flight was Cancelled Flightled...
14634      @AmericanAir right on cue with the delays
14636  @AmericanAir leaving over 20 minutes Late Flig...
14638  @AmericanAir you have my money, you change my ...
```

Name: text, Length: 9157, dtype: object

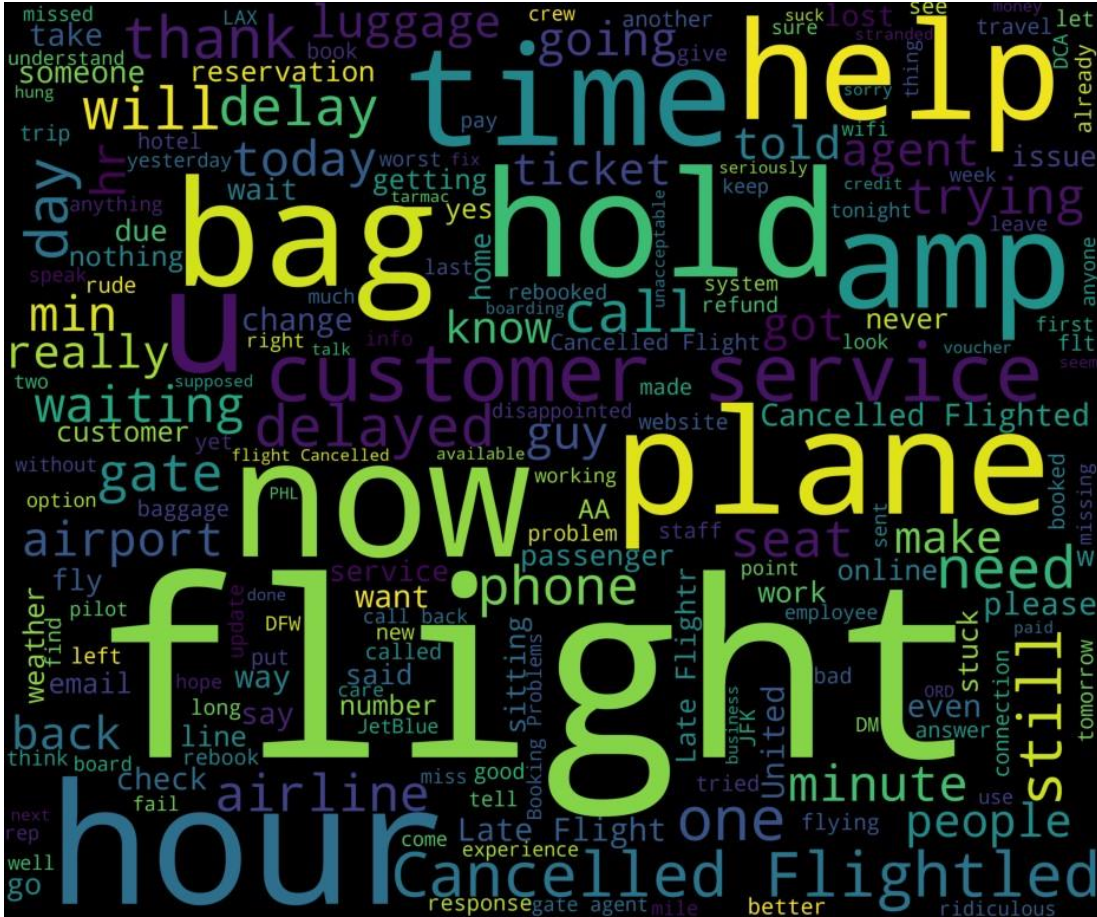
```
[34]: count_vect = CountVectorizer(stop_words="english")
neg_matrix = count_vect.
    ↪ fit_transform(data[data["airline_sentiment"]=="negative"]["text"])
freqs = zip(count_vect.get_feature_names_out(), neg_matrix.sum(axis=0)).
    ↪ tolist()[0])
# Sort from largest to smallest
print(sorted(freqs, key=lambda x: -x[1])[:100])
```

```

[('flight', 2937), ('united', 2899), ('usairways', 2375), ('americanair', 2089),
('southwestair', 1214), ('jetblue', 1051), ('cancelled', 921), ('service', 746),
('hours', 646), ('just', 622), ('help', 618), ('hold', 611), ('customer', 609),
('time', 596), ('plane', 530), ('delayed', 505), ('amp', 503), ('hour', 452),
('flightled', 445), ('http', 436), ('flights', 419), ('bag', 415), ('gate',
410), ('ve', 398), ('don', 388), ('late', 377), ('need', 373), ('phone', 367),
('waiting', 341), ('thanks', 315), ('got', 298), ('airline', 294), ('like',
291), ('trying', 288), ('delay', 272), ('wait', 272), ('today', 269),
('minutes', 266), ('day', 251), ('going', 249), ('bags', 245), ('luggage', 245),
('told', 245), ('airport', 244), ('people', 242), ('worst', 241), ('fly', 237),
```

('really', 236), ('did', 227), ('guys', 224), ('weather', 224), ('lost', 221), ('agent', 218), ('hrs', 217), ('way', 212), ('make', 211), ('change', 210), ('seat', 208), ('flighted', 205), ('want', 205), ('check', 204), ('know', 201), ('days', 200), ('home', 194), ('virginamerica', 191), ('baggage', 190), ('getting', 181), ('sitting', 179), ('ticket', 176), ('tomorrow', 176), ('let', 174), ('min', 171), ('customers', 169), ('flying', 168), ('line', 164), ('email', 163), ('online', 163), ('experience', 162), ('didn', 161), ('stuck', 160), ('work', 159), ('bad', 157), ('number', 156), ('won', 156), ('said', 155), ('seats', 154), ('30', 153), ('10', 150), ('problems', 150), ('times', 150), ('crew', 149), ('flightr', 148), ('doesn', 146), ('good', 145), ('ll', 144), ('aa', 143), ('travel', 142), ('yes', 142), ('response', 139), ('miss', 137)]

```
[35]: new_df = data[data["airline_sentiment"] == "positive"]
      words = " ".join(new_df["text"])
      cleaned_word = " ".join([word for word in words.split() if "http" not in word_
      ↪and not word.startswith("@") and word != "RT"]])
      wordcloud = WordCloud(stopwords = STOPWORDS,
      ↪background_color = "black", width = 3000, height = 2500).
      ↪generate(cleaned_word)
      plt.figure(figsize = (12, 12))
      plt.imshow(wordcloud)
      plt.axis("off")
      plt.show()
```

```
[37]: data.drop(data.loc[data["airline_sentiment"] == "neutral"].index, inplace =
      ↪ True)
```

```
[38]: from sklearn.preprocessing import LabelEncoderle
      = LabelEncoder()

      le.fit(data["airline_sentiment"])
      data["airline_sentiment_encoded"] = le.transform(data["airline_sentiment"])
      data.head()
```

```
[38]: airline_sentiment      text \
1      positive  @VirginAmerica plus you've added commercials t...
3      negative  @VirginAmerica it's really aggressive to blast...
4      negative  @VirginAmerica and it's a really big bad thing...
5      negative  @VirginAmerica seriously would pay $30 a fligh...
6      positive  @VirginAmerica yes, nearly every time I fly VX...
```

airline sentiment encoded

1	1
3	0
4	0
5	0
6	1

```
[39]: def tweet_to_words(tweet):
    letters_only = re.sub("[^a-zA-Z]", " ", tweet)
    words = letters_only.lower().split()
    stops = set(stopwords.words("english"))
    meaningful_words = [w for w in words if not w in stops]
    return " ".join(meaningful_words)
```

```
[40]: nltk.download("stopwords")
data["clean_tweet"] = data["text"].apply(lambda x: tweet_to_words(x))
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Administrator\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
[41]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11510 entries, 1 to 14638
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   airline_sentiment                    11510 non-null  object
1   text                                11510 non-null  object
2   airline_sentiment_encoded            11510 non-null  int32
3   clean_tweet                          11510 non-null  object
dtypes: int32(1), object(3)
memory usage: 404.6+ KB
```

```
[42]: X = data["clean_tweet"]
      y = data["airline_sentiment"]
```

```
[43]: print(X.shape, y.shape)
```

```
(11510,) (11510,)
```

```
[44]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 42)
      print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(8632,) (2878,) (8632,) (2878,)
```

```
[45]: vect = CountVectorizer()
      vect.fit(X_train)
```

[45]: CountVectorizer()

```
[46]: X_train_dtm = vect.transform(X_train)
      X_test_dtm = vect.transform(X_test)
```

```
[47]: vect_tunned = CountVectorizer(stop_words = "english", ngram_range = (1, 2),
      ↪ min_df = 0.1, max_df = 0.7, max_features = 100)
      vect_tunned
```

[47]: CountVectorizer(max_df=0.7, max_features=100, min_df=0.1, ngram_range=(1, 2), stop_words='english')

```
[48]: from sklearn.svm import SVC
      model = SVC(kernel = "linear", random_state = 10)
      model.fit(X_train_dtm, y_train)
      pred = model.predict(X_test_dtm)
      print("Accuracy Score: ", accuracy_score(y_test, pred) * 100)
```


Accuracy Score: 90.7574704656011

```
[49]: print("Confusion Matrix\n\n", confusion_matrix(y_test, pred))
```

Confusion Matrix

```
[[2179  122]
 [ 144  433]]
```

```
[50]: #defining the size of the canvas
      plt.rcParams["figure.figsize"] = [15,8]
      #confusion matrix to DataFrame
      conf_matrix = pd.DataFrame(data = confusion_matrix(y_test, pred), columns =
      ↪ ["Predicted:0", "Predicted:1"], index = ["Actual:0", "Actual:1"])
      #plotting the confusion matrix
      sns.heatmap(conf_matrix, annot = True, fmt = "d", cmap = "Paired", cbar =
      ↪ False, linewidths = 0.1, annot_kws = {"size":25})
      plt.xticks(fontsize = 20)
      plt.yticks(fontsize = 20)
      plt.show()
```



```
[51]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
negative	0.94	0.95	0.94	2301
positive	0.78	0.75	0.77	577
accuracy			0.91	2878
macro avg	0.86	0.85	0.85	2878
weighted avg	0.91	0.91	0.91	2878

```
[ ]:
```